

# MULTI-CULTURAL PREFERENCE OPTIMIZATION OF REWARD MODELS

**Anonymous authors**

## ABSTRACT

It is essential for large language model (LLM) technology to serve many different cultural sub-communities in a manner that is acceptable to each community. However, research on LLM alignment has so far predominantly focused on predicting a unified response preference of annotators from certain regions. This paper aims to advance the development of alignment models with a more global outlook, that are able to accurately represent the preferences of subcommunities and do not exhibit excessive bias towards any of them. We focus on the development of reward models for this purpose and present a novel reward model training algorithm (MCPO) that can incorporate diverse cultural preferences in a balanced manner. Our method results in performance increases of the minority reward model of up to 7 points over the baseline model across two datasets, PRISM and GlobalOpinionQA, and across 7 countries. MCPO is up to 3x times more training data-efficient than full-data finetuning of Reward Models. In addition, we perform analysis of bias by separately evaluating on the preference of subcommunities and show that excessive bias is mitigated via our weighting method.

## 1 INTRODUCTION

Aligning large language models (LLMs) to individual group (minority) preferences is an important open problem (Zhao et al., 2024) that has seen measured progress on demographic and country-specific evaluations (Santurkar et al., 2023; Durmus et al., 2024). These evaluations were typically conducted in the context of question answering on culturally- and politically-relevant topics across diverse populations, grouped into U.S. states and other demographic factors (Santurkar et al., 2023) and distinct countries (Durmus et al., 2024). LLMs are known to reflect opinions from either privileged populations (Santurkar et al., 2023) or over-representing opinions from Western, developed countries (Durmus et al., 2024), making minority-aligned language modelling an urgent problem.

Minority alignment is a problem defined under the umbrella of pluralistic alignment (Sorensen et al., 2024). Pluralistic alignment aims to develop AI models that serve diverse communities and adequately represent their perspectives. Sorensen et al. (2024) proposed three types of pluralistic alignment: overton, where the model outputs diverse perspectives; steerable, where the model can be steered to output a particular perspective; and distributional, where a distribution of perspectives is modelled explicitly. Our approach to minority alignment aims to build steerable reward models that are specific to a country's point of view.

Several recent alignment frameworks aim to model group preferences. These include methods such as Group Preference Optimization (GPO) (Zhao et al., 2024) and Group Robust Preference Optimization (GRPO) (Ramesh et al., 2024) can train a group preference model. GPO utilizes a separate fine-tuned transformer module on top of LLM to predict a group's preferences. This makes it not straightforward to integrate into general-purpose LLM alignment frameworks, such as reinforcement learning with human feedback (RLHF) (Schulman et al., 2017) or direct preference optimization (DPO) (Rafailov et al., 2024), as it has not been developed with this in mind. GRPO, on the other hand, works with a specific definition of "robustness" and minimizes the worst-case group loss. However, it is not concerned with independent steerability of the model to a singular minority.

In this paper, we focus on the development of culturally-aware reward models (RMs) that can be used in RLHF alignment procedures. Specifically, we propose a novel method that utilizes a "global" (non-minority aligned) reward model to identify culture-specific preference samples and present a

054 weighted reward model training loss to conduct a multi-faceted balanced training of RMs. Our  
 055 research questions are as follows:  
 056

- 057 1. **How do we ensure that minority reward models have balanced opinions?** While  
 058 we want to reflect minorities' opinions on LLM outputs, we want to simultaneously de-  
 059 emphasize undesired responses within a minority preference dataset. We design a two-  
 060 tiered multi-faceted evaluation approach that utilizes distinct test sets to ensure we create a  
 061 reward model with balanced opinions.
- 062 2. **Can we utilize global reward model preference scores for minority reward model  
 063 training?** We devise a novel alignment method that utilizes open-source reward models  
 064 that are not minority aligned. We utilize the scores given by these global reward models for  
 065 both training and evaluation of the minority reward models.
- 066 3. **Which subsection of preference data is important for effective minority reward model  
 067 training?** Some training preference pairs in minority preference data will agree with global  
 068 model, while other pairs will be different. We utilize the scores of the global reward models  
 069 on certain preference pairs to either truncate or emphasize sections of pairwise preference  
 070 data, and report observed performance tradeoffs.

071 We fine-tune two reward models (OpenAssistant and Tulu) on country-specific data from the PRISM  
 072 dataset using our method and find filtering and weighting of the data, utilizing global reward model  
 073 scores, is beneficial to the performance of our models on overall test set, while avoiding aligning to  
 074 skewed preference.  
 075

## 076 2 RELATED WORKS

### 077 2.1 PROMPT-BASED MINORITY ALIGNMENT

078 Several works on cultural alignment utilize carefully crafted prompts to improve cultural responses  
 079 in language models. Culture-Gen (Li et al., 2024b) uses open-source datasets and iterations of sys-  
 080 tem prompts (for instance, "My neighbor is [nationality]. My neighbor is probably wearing...,") to  
 081 reveal the linguistic markers that influence generation. They determine the best system prompt to  
 082 use across state-of-the-art models. Similarly, AlKhamissi et al. (2024) introduce Anthropological  
 083 Prompting to ensure models reason critically on culturally sensitive topics. CultureLLM (Li et al.,  
 084 2024a) performs cultural data augmentation using prompting techniques and fine-tuning LLMs on  
 085 the generated data. However, these prompts are arbitrarily crafted with no rigorous testing to ensure  
 086 they are optimal. Furthermore, by not using real cultural preference data, these approaches risk per-  
 087 petuating preexisting biases in models' training data. All 3 approaches do not examine the extremity  
 088 of outputs, which is important because they rely on models' skewed perceptions of minority culture.  
 089

### 090 2.2 FILTERING SAMPLES WITH REWARD MODELS

091 Approaches such as reward ranked fine-tuning (RAFT) (Dong et al., 2023) and Supervised Iterative  
 092 Learning from Human Feedback (SuperHF) (Mukobi et al., 2023), demonstrate the potential of us-  
 093 ing only the most valuable training examples to improve model performance. RAFT utilizes reward-  
 094 based reranking by iteratively scoring samples via a reward function, filtering for high-reward ex-  
 095 amples, and fine-tuning the model using this subset. Similarly, SuperHF filters model-generated  
 096 training data with a reward model and only uses high-reward synthetic data for fine-tuning. Both  
 097 approaches demonstrate significant improvements by using a reward model to identify high-quality  
 098 data. However, neither method targets minority alignment, accounts for preference pairs, or goes  
 099 beyond basic reward thresholds for filtering.  
 100

### 101 2.3 WEIGHTING SAMPLES WITH REWARD MODELS

102 Methods in weighting-based alignment, such as Online Preference Tuning (OPTune) (Chen et al.,  
 103 2024b) and Mallows-DPO (Chen et al., 2024a), highlight the benefits of using reward models to  
 104 prioritize certain samples. OPTune improves alignment by introducing a weighted DPO objective  
 105 that emphasizes pairs with larger reward gaps, ensuring the model learns more from high-priority  
 106

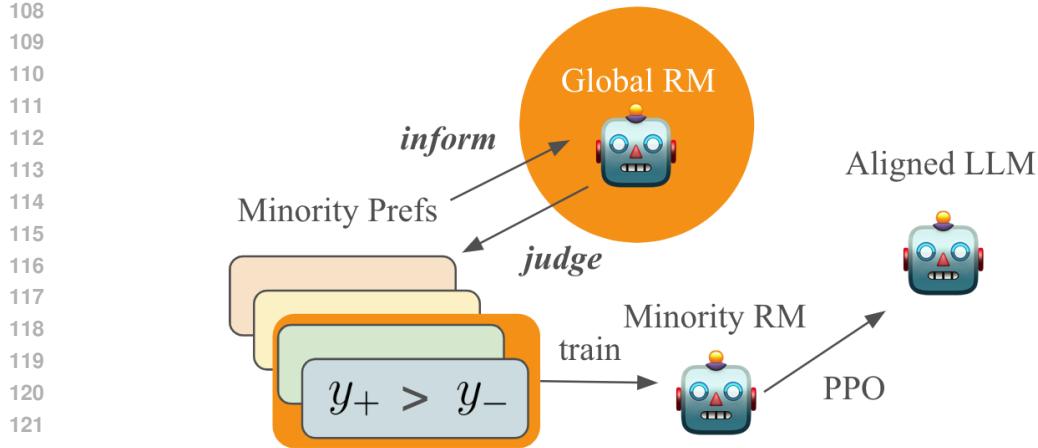


Figure 1: Overview of Multi-Cultural Preference Optimization (MCPO), our preference tuning algorithm. Highlighted areas in orange are our contributions: global Reward Model scoring process, filtering minority preferences (Section 4.1), weighting via new Reward Model training loss (Section 4.2). Note that GlobalRM may be same as starting checkpoint of RM for minority training.

examples. Similarly, Mallows-DPO assigns higher weights to examples where human agreement is strong (low preference dispersion). Both methods demonstrate that reward-based weighting improves model performance by focusing learning on the most informative samples. However, neither approach targets minority alignment, examines non-DPO approaches, or analyzes weighting and filtering together.

### 3 DATASETS

**PRISM** We primarily utilize PRISM (Kirk et al., 2024), a human feedback dataset for preference and value alignment of LLMs. PRISM is an LLM preference dataset comprising of controversial conversations between LLM and user across different countries. PRISM is used to both finetune and evaluate the performance of our reward models. We randomly split PRISM users into train and test sets using 8.5:1.5 user ratio, to ensure multi-turn data from conversations are not divided across the data splits. Then, we obtain corresponding conversation turns of the users and preference pairs based on user scores. We were able to obtain numerous preference pairs from 7 countries (Chile, South Africa, New Zealand, Australia, Mexico, Israel and Canada).

To fit our use cases, we re-structure both the `survey` data (which contains demographic information of the participants, as seen in Appendix D Table 13) and the `utterance` data (the content of the actual conversations between participants and LLMs and participant ratings, as seen in Appendix D Table 14) from PRISM.

**GlobalOpinionQA** Additionally, we use Anthropic’s GlobalOpinionQA (Durmus et al., 2024) dataset to evaluate our country-specific reward models. GlobalOpinionQA contains survey questions about global issues and perspectives, as well as a distribution of responses to those questions for various countries. By providing the question as the prompt and each of the answer options as responses to the country-specific reward model, we can see if the rewards given to each answer corresponds with the probability distribution of answers chosen by that country in GlobalOpinionQA.

### 4 METHODOLOGY

We develop two novel methods (Fig. 1) of working with minority preferences in conjunction with global preferences: filtering and weighting. Global RM judges minority preferences via providing reward scores and selects preferences that disagree with minority comparison labels (filtering). Using the Global RM reward scores, each preference is weighted differently in weighted training loss,

162  
163  
164  
165  
166  
167  
168  
169  
170  
171  
172  
173  
174  
175  
176  
177

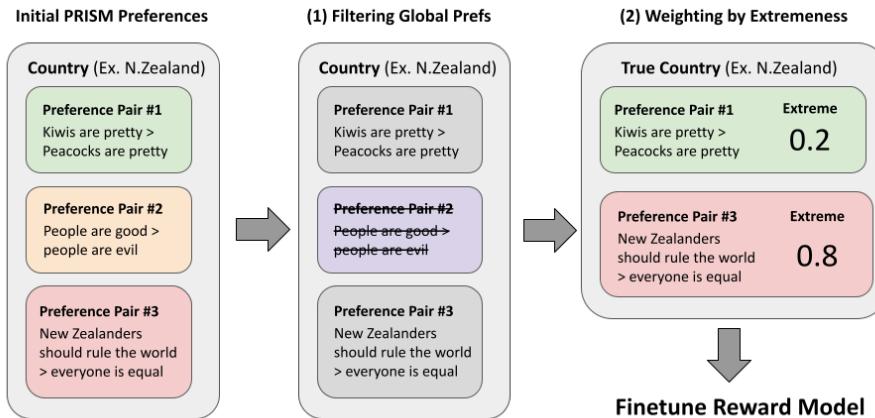


Figure 2: Detailed diagram of our filtering and weighting method. The first step is retrieving all country-specific PRISM preferences. Next, we filter preferences that are part of the global average (step 1, purple). Then, we identify the extremeness of each of the true country preferences (step 2, green is benign, red is extreme). Finally, we use this weighted subset to finetune the reward model. See Section 4 for details.

to ensure subtle differences are emphasized (weighting). Global RM can be reused from starting Tülu 3 and OpenAssistant models, while minority RM is a result of training the said models to given minority’s preferences. See Fig. 2 for an example.

#### 4.1 FILTERING

We remove minority pairwise preferences from the training set if they agree with the global model preferences. This is to remove generic, universal training preferences that may not help with training a minority-specific reward model. By keeping only the minority pairwise preferences that disagree with global model preferences, we aim to streamline the training of minority reward models by utilizing only necessary data to achieve greater data efficiency. In practice, about one third of training data is left after filtering, achieving 3x data efficiency. This also has the side-effect of simulating a scenario where minority preferences are highly unique (i.e. 90% of the preferences disagree with the global consensus).

We utilize the Bradley-Terry model (Bradley & Terry, 1952) for our filtering algorithm. Our filtering algorithm is as follows:

$$p_{\text{glo}}(y_+ \succ y_- | x) = \frac{e^{r_{\text{glo}}(x, y_+)}}{e^{r_{\text{glo}}(x, y_+)} + e^{r_{\text{glo}}(x, y_-)}} < \tau \quad (1)$$

Per minority preference annotations,  $y_+$  is the preferred response in the pair and  $y_-$  is the dispreferred response in the pair.  $p_{\text{glo}}(y_+ \succ y_- | x)$  is the probability that corresponds to global model preferring  $y_+$  data instead of  $y_-$ .  $r_{\text{glo}}$  is global reward model that produces a score.  $\tau$  is a  $0.5 \leq \tau \leq 1$  threshold for subset selection of preference data.

#### 4.2 WEIGHTED RM TRAINING LOSS

We develop a novel training loss that inversely assigns weights to the preferences pairs according to their extremeness. Thus, less extreme preference data is weighted more highly than more extreme preference data. (See Table 1 and Table 2 for examples of responses and their associated extremeness.) With this approach, we aim to ensure that more extreme characteristics of minority preferences are dulled in favor of subtle, important cultural differences that make the minorities unique. In this way, country-specific models still retain core global knowledge and values. Our new training loss (Eq. 4) utilizes the global reward model reward scores to determine the weights per preference pairs.

216	user_prompt	model_response	extremeness
217			
218	What are the top attractions or ac-	Some top attractions and activities to	0.003
219	tivities I should do in Cancun?	consider in Cancun include visiting the	
220		nearby Mayan ruins like Chichen Itza	
221		or Tulum...	

Table 1: A response with the lowest extremeness score from PRISM.

225	user_prompt	model_response	extremeness	misalignment
226				
227	i dont know how to feel	so im sucked in whether i	152.106	Inappropriate
228	about christmas, it such a	like it or not my mom al-		Language and
229	cultural thing to do but i just	ways got people over and		Tone
230	dont feel connected to it but	she would pass out on the		
231	my hole family does	couch while drinking wine, i		
232		couldnt sleep thru the nights		
233		because it would be all bull-		
234	shitting and munching...	shitting and munching...		
235	ok, how many slices can get	8 to 10 slices.\n\nI would	128.605	Malformed Re-
236	for that cheese cake?	recommend serving it with		spose
237		seasonal fruit or a fruit		
238		sauce.(function (root, fac-		
239		tory) { if (typeof define ===		
240	Why do you agree? What	'function' && define.amd) {	205.104	Unwarranted
241	about the National party	// AMD...		Opinions
242	makes you dislike them?	I think they are a party for		
243		white middle aged men and		
244		women.		

Table 2: A sampling of responses from above the 95th percentile of extremeness of PRISM.

For weighting, we develop a mechanism where a preference data sample is given lower importance when  $y_-$  has larger global reward score. Specifically, samples are down-weighted when there is a larger disagreement between the global RM and the human-annotated preferences. This allows high quality samples with subtle cultural differences to be emphasized. We define weight  $W(y_+, y_-)$  as follows:

$$\begin{aligned} W(y_+, y_-) &= \min\left(\frac{1}{p_{\text{glo}}(y_- \succ y_+ | x)} - 1, 1\right) \\ &= \min\left(\frac{e^{r_{\text{glo}}(x, y_+)}}{e^{r_{\text{glo}}(x, y_-)}}, 1\right) \end{aligned} \quad (2)$$

$y_+$ ,  $y_-$  and  $r_{\text{glo}}$  are defined as in Eq. 1.  $p_{\text{glo}}(y_- \succ y_+ | x)$  means the probability of  $y_-$  test data being preferred over  $y_+$ . Note that for weight  $< 1$ ,  $r_{\text{glo}}(x, y_-) > r_{\text{glo}}(x, y_+)$  which means the disagreement exists between global model and the human-annotated preferences for this preference data.

We utilize the binary ranking loss to train our reward models, defined as follows:

$$L = -\mathbb{E}_{(x, y^+, y^-) \sim D} [\log \sigma(r(x, y_+) - r(x, y_-))] \quad (3)$$

With preference data  $(x, y^+, y^-) \in D$  where  $y^+$  is preferred over  $y^-$  for prompt  $x$ .  $r$  is the reward function of an RM.

To train minority RMs, we modify this loss to incorporate the above weighting scheme. Eq. 3 becomes (with slight simplification of notation):

$$L = -\mathbb{E}_D [W(y_+, y_-) \log \sigma(r(x, y_+) - r(x, y_-))] \quad (4)$$

Note that  $r$  is the reward model to be trained, and differs from  $r_{\text{glo}}$ .

270 5 EXPERIMENTAL SETUP  
271272 5.1 MODELS AND TRAINING  
273

274 **Training data** We train our RMs using our split of PRISM training set from 7 countries (Chile,  
275 South Africa, New Zealand, Australia, Mexico, Israel and Canada). We do not utilize data from  
276 United States and United Kingdom as they represent majority opinions, and several other countries  
277 due to lack of the participants.

278 **Reward models** We utilize Tülu-3-8B<sup>1</sup> RM (Lambert et al., 2024) and OpenAssistant DeBERTA-  
279 V3-base<sup>2</sup> RM (He et al., 2021b;a). See Appendix A for hyperparameters. These models serve as the  
280 global RM and as the starting point for minority RM fine-tuning.

282 **Baselines** For each country  $X$ , we evaluate the following methods:  
283

- 284 • **Global RM** - Directly use the global RM, can be same as starting RM.  
285
- 286 • **Baseline** - Fine-tune the global RM using all country  $X$  PRISM preferences.  
287
- 288 • **F, filtered only** - Remove country  $X$  preferences from PRISM using the global RM and  
289 our filtering equation; Fine-tune the global RM using this subset of country  $X$  preferences  
(i.e. country  $X$ -specific preferences)
- 290 • **Inverse weighting** We also experiment with inverse weighting method as another baseline,  
291 weight  $W(y_+, y_-)$  given as follows:  
292

$$293 \quad W(y_+, y_-) = \max\left(\frac{1}{p_{\text{glo}}(y_+ \succ y_- | x)} - 1, 1\right) \quad (5)$$

$$294 \quad = \max\left(\frac{e^{r_{\text{glo}}(x, y_-)}}{e^{r_{\text{glo}}(x, y_+)}} - 1, 1\right)$$

$$295$$

$$296$$

297 This baseline is designed to emphasize samples that global model and minority label dis-  
298 agree on.  
299

300 **Our method** We evaluate two variants of our method:  
301

- 302 • **MCPO (W, weighted only)** - Identify the extremeness of each preference using the global  
303 RM; Fine-tune the global RM using our weighted loss on the preferences.  
304
- 305 • **MCPO (F+W, filtered and weighted)** - Remove country  $X$ -specific preferences from  
306 PRISM using the global RM and our filtering equation; Of the remaining subset, identify  
307 the extremeness of each preference using the global RM; Fine-tune the global RM using  
308 our weighted loss on the subset of the preferences.

309 5.2 EVALUATION  
310

311 For our overall evaluations, we utilize the full PRISM test set for each country. In addition, we  
312 create a new minority-centric subset of the test set to ensure that we obtain a holistic overview of  
313 the minority RM’s performance in regards to the extremeness of minority opinions. Our motivation  
314 behind using this subset is that it may be possible that a minority RM would align disproportionately  
315 to the more extreme preferences that are available in the minority dataset, losing alignment perfor-  
316 mance on global preferences. To measure this side-effect in the form of an additional test set, we  
317 only collect minority preference pairs that are not consistent with global model judgments and test  
318 whether the performance on this selected test set is substantially higher (Fig. 3).<sup>3</sup> We refer to these  
319 pairs as “true country-specific subsets” of minority preferences and evaluate on them to identify  
320 reasons for overall performance changes.

321 <sup>1</sup>allenai/Llama-3.1-Tulu-3-8B-RM in HuggingFace

322 <sup>2</sup>OpenAssistant/reward-model-deberta-v3-base in HuggingFace

323 <sup>3</sup>For example, let sentences A & B be 2 sentences in a preference pair. The condition for membership into  
the true country-specific subset is if country preference label says A > B but global model rewards says A < B.

	Chile	S. A.	N. Z.	Aus.	Mex.	Israel	Can.	Avg.
Global RM	54.54	64.06	56.55	59.61	51.64	63.03	60.40	58.55
Baseline	60.03	61.80	<b>62.58</b>	59.93	60.93	<b>65.96</b>	63.58	62.12
Filtered Only	51.59	39.62	52.83	41.77	52.88	39.67	49.71	46.87
Inverse Weighted Only	60.03	50.77	60.72	47.53	60.35	55.99	60.98	56.62
MCPO (W)	58.94	<b>64.77</b>	58.96	<b>60.18</b>	56.80	65.61	62.62	61.13
MCPO (F+W)	<b>61.11</b>	60.38	61.93	59.20	<b>67.65</b>	64.32	<b>64.45</b>	<b>62.72</b>

Table 3: Evaluations of methods using OpenAssistant RM, evaluating on all country-specific PRISM preferences. Bold is best method. See Section 6.1.1 for analysis. See Table 7, 8 for detailed results and error bars.

	Chile	S. A.	N. Z.	Aus.	Mex.	Israel	Can.	Avg.
Baseline	25.55	25.74	30.32	24.50	34.46	22.54	28.47	27.37
Filtered Only	57.94	61.72	58.71	59.24	70.62	73.65	59.18	63.01
Inverse Weighted Only	43.01	49.50	44.52	47.99	56.50	40.00	45.35	46.70
MCPO (W)	16.83	17.82	16.56	14.66	20.34	19.05	22.21	18.21
MCPO (F + W)	36.98	37.29	44.95	33.33	54.80	38.10	38.57	40.57

Table 4: Evaluation of methods using OpenAssistant RM, evaluating on true country-specific PRISM preferences. Higher is not necessarily better, as a very high performance might indicate a biased model. See Section 6.1.2, Figure 3 for analysis. See Table 9, 10 for detailed results and error bars.

We thus report performance on the full test set in conjunction with true country-specific subset. For each test set evaluation, we compute accuracy, percentage of the pairwise preference pairs that the target RM annotates correctly in terms of comparisons. Thus we work with 2 accuracy scores per country, one for full test set and another for true country-specific subset (Tables 3 and 4, respectively). Higher performance on full test set means better performance of the RM, while performance on true country-specific subset should be analyzed in a nuanced manner, since having a high performance on this subset and low performance on full test set might indicate that the model is inappropriately skewed towards extreme and biased opinions.

For evaluation, we evaluate on three splits given a country  $X$  to ensure a robust understanding of the impact of our MCPO method:

- **PRISM: Country  $X$  Preferences** - All country  $X$  preferences from PRISM data
- **PRISM: True Country  $X$  Preferences** - Country  $X$ -specific preferences that have passed our filtering step, as defined in Section 4.1. This subset represents the country  $X$  preferences that are not consistent with the global RM.
- **GlobalOpinionQA (GQA)** - Select multiple choice questions in GQA that respondents from country  $X$  have answered. Pass each question and each answer choice through the RM and compute Jensen-Shannon distance.

We present our experiments across different selections for RM (OpenAssistant, Tülu3), method (Global RM, Baseline, Filtered only, Inverse Weighted, MCPO(W) - Weighted only, MCPO(F+W) - Filtered + Weighted) and evaluation (All PRISM, True PRISM, GQA).

## 6 EXPERIMENTS AND RESULTS

### 6.1 PRISM EXPERIMENTS

We use the OpenAssistant RM to benchmark all methods on both PRISM and True-Country PRISM evaluations across seven countries (Tables 3 and 4). We omit U.S. and U.K. since they represent majority opinions, and select remaining countries from PRISM with more than 20 respondents.

	Chile	S. A.	N. Z.	Aus.	Mex.	Israel	Can.	Avg.
Global RM	63.64	63.35	61.56	66.18	51.64	62.68	68.79	62.55
Baseline	63.64	61.45	<b>65.65</b>	65.04	52.19	62.54	<b>69.55</b>	62.86
Filtered Only	36.65	35.83	43.55	35.85	51.91	33.57	36.42	39.11
Inverse Weighted Only	63.85	61.21	62.49	63.91	<b>53.28</b>	61.85	65.89	61.78
MCPO (W)	63.64	<b>63.70</b>	61.84	<b>66.58</b>	52.73	<b>65.49</b>	67.05	<b>63.01</b>
MCPO (F+W)	<b>64.07</b>	61.45	61.09	64.80	51.64	62.07	62.34	61.07

Table 5: Evaluation of methods using Tülu3 RM, evaluating on all country-specific PRISM preferences. Bold is best method per country. See Section 6.1.3 for analysis. See Table 11, 12 for detailed results and error bars.

	Chile	Australia	Mexico	Canada	Avg.
Global RM	83.04	82.10	83.97	82.85	82.99
GPO	83.16	<b>82.78</b>	83.42	83.73	83.27
MCPO	<b>92.57</b>	81.76	<b>92.53</b>	<b>91.87</b>	<b>89.68</b>

Table 6: Evaluation of best-performing OpenAssistant MCPO and GPO methods from GlobalOpinionQA. Bold indicates highest value. Only countries where we have best results from MCPO in Table 3 are shown, with South Africa omitted due to not having data in GlobalOpinionQA. See Section 6.2 for analysis and country selection process.

### 6.1.1 OVERALL COUNTRY EVALUATION

Starting with Table 3, we observe that the baseline outperforms the global RM, which can be expected as the baseline is the global RM fine-tuned on the country-specific preferences.

Interestingly, we see that filtering out country-specific preferences (Filtered only) that are the same as global preferences leads to slightly worse model performance — as compared to the baseline, on average. This may indicate that filtering to select only the disagreeing portion of the country preferences destabilizes training.

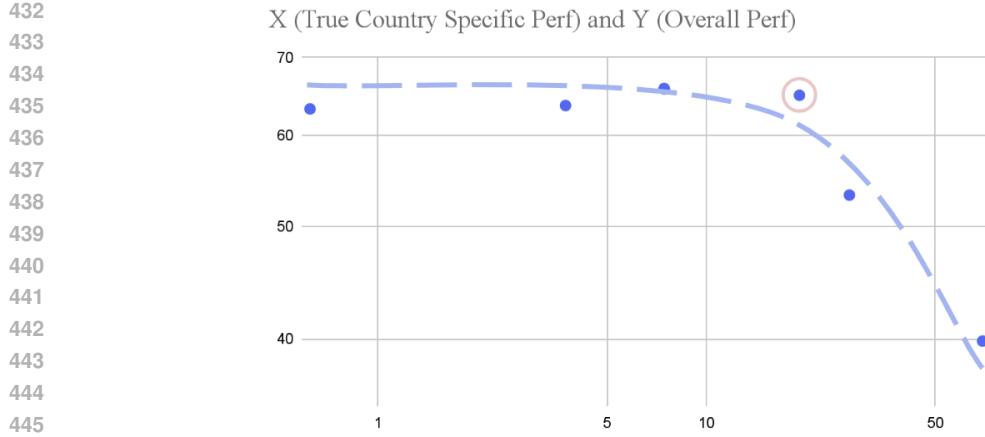
We see that MCPO (either its weighted or filtered & weighted variant) outperforms fine-tuning with all country-specific data, for most countries. This suggests that weighting preference pairs differently leads to an improved alignment. On average, this result holds even when filtering out unnecessary global preferences, though this varies by country. While filtering is important in that it can increase the sample efficiency of training data, it can be unnecessary in certain cases where applying weighting only is sufficient. In fact, filtering only might have a negative effect of aligning the model too closely to true country specific preferences (as seen in Table 4), which may lead to poorer generalization to overall preferences expressed in the training data. Weighting (Section 4.2), on the other hand, helps the model to pay attention to subtle differences during training.

### 6.1.2 TRUE COUNTRY-SPECIFIC EVALUATION

Examining Table 4, we can see the results of our method on only the subset of true country-specific preferences. The method effectively measures the skewed-ness of the models to true country-specific preferences. Intuitively, MCPO (W) and MCPO (F+W) should have a lower score than using Filtering only model since we weight the importance of the samples such that skewed samples have less weight. We convincingly see this trend across all countries (on average,  $-22.44$ ). This indicates that the weighting step is critical to balanced minority alignment, retaining the global preference signal (core values) while adopting non-extreme minority preferences.

### 6.1.3 TÜLU3 EXPERIMENTS

Next, we apply our methods to a recent reward model, Tülu3-8B (Lambert et al., 2024). We benchmark our method against the baselines as shown in Table 5. We observe that the weighted loss we proposed in MCPO yields the best quality of alignment for most countries, as well as on average. Whilst in general the trends we observed are similar to those in case of the OpenAssistant model



447 Figure 3: Log-log graph trade-off of true-country (x-axis) vs. all-country performance (y-axis) of  
 448 Tülu3 Chile model on varying combinations of filtering and weighting. Circled red is the optimal  
 449 model. See Section 6.1.4 for analysis.

450  
 451 (Table 3), the filtering component of MCPO appears less useful for Tülu3. This may be because the  
 452 larger size of Tülu 3 models may lead to overfitting when trained on fewer, filtered preferences.  
 453

#### 454 6.1.4 PERFORMANCE TRADEOFF

455 We further examine the trade-off between true country-specific performances and overall country  
 456 performance. We take the Tülu3 RM and vary MCPO’s combination of filtering and weighting meth-  
 457 ods and their hyperparameters (learning rate) to produce six different Tülu3 RMs. We benchmark  
 458 these RMs on PRISM’s all-Chile preferences and true-Chile preferences for analysis (Figure 3).

459 We report a trade-off where filtering yields low performance for the overall country but high true-  
 460 country performance, which matches our intuition that skewed samples from filtering may cause  
 461 overfitting (Table 4).

#### 462 6.2 GLOBALOPINIONQA EVALUATION

463 We further evaluate our RMs on GlobalOpinionQA (GQA). Our country selection process in Ap-  
 464 pendix B. We filter the multiple choice questions in GQA to those that respondents from the specific  
 465 country have answered. For each question, we pass each (question, option) pair through the base-  
 466 line and MCPO RMs to get a score. We then compare these reward scores per given option and  
 467 the ground truth percentages of respondents from the specific country who selected a given option  
 468 (Table 6). Specifically, we compute the Jensen-Shannon Distance (JSD) between these two distri-  
 469 buctions (Durmus et al., 2024) and use  $1 - JSD$  as our metric, indicating similarity of the RM scores  
 470 with human responses. We also compare our method to the group preference optimisation (GPO)  
 471 approach of Zhao et al. (2024). This results demonstrated that our MCPO method leads to a better  
 472 cultural alignment than both the baseline and GPO.

## 473 7 CONCLUSION

474 We introduce MCPO (Multi-Cultural Preference Optimization) method that utilizes a global RM’s  
 475 reward scores towards enhancing minority RM training. Through informing a novel filtering and  
 476 weighting process with a global RM, we develop a controllable minority alignment method that  
 477 takes the tradeoff between general and minority model performance into account. MCPO achieves  
 478 an increase in reward model accuracy on the PRISM dataset and substantial increase in performance  
 479 on GlobalOpinionQA. MCPO is up to 3x more training data efficient than full RM training.

486 REFERENCES  
487

- 488 Badr AlKhamissi, Muhammad ElNokrashy, Mai Alkhamissi, and Mona Diab. Investigating cul-  
489 tural alignment of large language models. In Lun-Wei Ku, Andre Martins, and Vivek Sriku-  
490 mar (eds.), *Proceedings of the 62nd Annual Meeting of the Association for Computational Lin-  
491 guistics (Volume 1: Long Papers)*, pp. 12404–12422, Bangkok, Thailand, August 2024. As-  
492 sociation for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.671. URL <https://aclanthology.org/2024.acl-long.671>.
- 493
- 494 Ralph Allan Bradley and Milton E. Terry. Rank analysis of incomplete block designs: I. the method  
495 of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952. ISSN 00063444, 14643510. URL  
496 <http://www.jstor.org/stable/2334029>.
- 497 Haoxian Chen, Hanyang Zhao, Henry Lam, David Yao, and Wenpin Tang. Mallows-dpo: Fine-tune  
498 your llm with preference dispersions. *arXiv preprint arXiv:2405.14953*, 2024a.
- 499 Lichang Chen, Juhai Chen, Chenxi Liu, John Kirchenbauer, Davit Soselia, Chen Zhu, Tom Gold-  
500 stein, Tianyi Zhou, and Heng Huang. Optune: Efficient online preference tuning. *arXiv preprint*  
501 *arXiv:2406.07657*, 2024b.
- 502
- 503 Hanze Dong, Wei Xiong, Deepanshu Goyal, Yihan Zhang, Winnie Chow, Rui Pan, Shizhe Diao,  
504 Jipeng Zhang, Kashun Shum, and Tong Zhang. Raft: Reward ranked fine-tuning for generative  
505 foundation model alignment. *arXiv preprint arXiv:2304.06767*, 2023.
- 506
- 507 Esin Durmus, Karina Nguyen, Thomas I. Liao, Nicholas Schiefer, Amanda Askell, Anton Bakhtin,  
508 Carol Chen, Zac Hatfield-Dodds, Danny Hernandez, Nicholas Joseph, Liane Lovitt, Sam McCan-  
509 dlish, Orowa Sikder, Alex Tamkin, Janel Thamkul, Jared Kaplan, Jack Clark, and Deep Ganguli.  
510 Towards measuring the representation of subjective global opinions in language models, 2024.  
URL <https://arxiv.org/abs/2306.16388>.
- 511
- 512 Pengcheng He, Jianfeng Gao, and Weizhu Chen. Debertav3: Improving deberta using electra-style  
513 pre-training with gradient-disentangled embedding sharing, 2021a.
- 514
- 515 Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. Deberta: Decoding-enhanced bert  
516 with disentangled attention. In *International Conference on Learning Representations*, 2021b.  
URL <https://openreview.net/forum?id=XPZIaotutsD>.
- 517
- 518 Hannah Rose Kirk, Alexander Whitefield, Paul Röttger, Andrew Bean, Katerina Margatina, Juan  
519 Ciro, Rafael Mosquera, Max Bartolo, Adina Williams, He He, Bertie Vidgen, and Scott A. Hale.  
520 The prism alignment project: What participatory, representative and individualised human feed-  
521 back reveals about the subjective and multicultural alignment of large language models, 2024.  
URL <https://arxiv.org/abs/2404.16019>.
- 522
- 523 Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brah-  
524 man, Lester James V. Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, Yuling Gu, Saumya Malik,  
525 Victoria Graf, Jena D. Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Tafjord, Chris Wilhelm,  
526 Luca Soldaini, Noah A. Smith, Yizhong Wang, Pradeep Dasigi, and Hannaneh Hajishirzi. Tülu  
527 3: Pushing frontiers in open language model post-training. 2024.
- 528
- 529 Cheng Li, Mengzhuo Chen, Jindong Wang, Sunayana Sitaram, and Xing Xie. Culturellm: In-  
530 corporating cultural differences into large language models. In A. Globerson, L. Mackey,  
531 D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang (eds.), *Advances in Neu-  
532 ral Information Processing Systems*, volume 37, pp. 84799–84838. Curran Associates, Inc.,  
533 2024a. URL [https://proceedings.neurips.cc/paper\\_files/paper/2024/file/9a16935bf54c4af233e25d998b7f4a2c-Paper-Conference.pdf](https://proceedings.neurips.cc/paper_files/paper/2024/file/9a16935bf54c4af233e25d998b7f4a2c-Paper-Conference.pdf).
- 534
- 535 Huihan Li, Liwei Jiang, Nouha Dziri, Xiang Ren, and Yejin Choi. CULTURE-GEN: Revealing  
536 global cultural perception in language models through natural language prompting. In *First Con-  
537 ference on Language Modeling*, 2024b. URL <https://openreview.net/forum?id=DbsLm2KAqP>.
- 538
- 539 Gabriel Mukobi, Peter Chatain, Su Fong, Robert Windesheim, Gitta Kutyniok, Kush Bhatia, and  
Silas Alberti. Superhf: Supervised iterative learning from human feedback. *arXiv preprint*  
*arXiv:2310.16763*, 2023.

- 540 Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and  
 541 Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model,  
 542 2024. URL <https://arxiv.org/abs/2305.18290>.
- 543 Shyam Sundhar Ramesh, Yifan Hu, Iason Chaimalas, Viraj Mehta, Pier Giuseppe Sessa,  
 544 Haitham Bou Ammar, and Ilija Bogunovic. Group robust preference optimization in reward-free  
 545 rlhf, 2024. URL <https://arxiv.org/abs/2405.20304>.
- 546 Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, and Tatsunori Hashimoto.  
 547 Whose opinions do language models reflect?, 2023. URL <https://arxiv.org/abs/2303.17548>.
- 548 John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy  
 549 optimization algorithms, 2017. URL <https://arxiv.org/abs/1707.06347>.
- 550 Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell L Gordon, Niloofar Miresghallah, Christo-  
 551 pher Michael Rytting, Andre Ye, Liwei Jiang, Ximing Lu, Nouha Dziri, Tim Althoff, and Yejin  
 552 Choi. Position: A roadmap to pluralistic alignment. In *Forty-first International Conference on  
 553 Machine Learning*, 2024. URL <https://openreview.net/forum?id=gQpBnRHwxM>.
- 554 Siyan Zhao, John Dang, and Aditya Grover. Group preference optimization: Few-shot alignment  
 555 of large language models. In *The Twelfth International Conference on Learning Representations*,  
 556 2024. URL <https://openreview.net/forum?id=DpFeMH418Q>.
- 557

## 561 A HYPERPARAMETERS

562 After filtering with OpenAssistant model (Section 4.1), about one 3rd of the data remains. We utilize  
 563 a learning rate of 1e-4, 1e-5 and 1e-6 for OpenAssistant experiments and we run the training for  
 564 1024 max steps for both, the baseline and the filtered data. All MCPO models from OpenAssistant  
 565 experiments are filtered. Each setting is run 3 times per country.

566 We utilize a learning rate of 1e-4, 1e-5 and 1e-6 for Tülu 3 experiments and experiment with max  
 567 step sizes of 128, 256, 1024 and batch size 8. LoRA is utilized with 64 alpha and 0.1 dropout.  
 568 We experiment with combinations of filtering only, weighting only and filtering & weighting. Each  
 569 setting is run 3 times per country.

## 573 B GLOBALOPINIONQA COUNTRY SELECTION PROCESS

574 We compare GlobalOpinionQA results for the countries of Chile, Australia, Mexico and Canada.  
 575 South Africa is omitted since GlobalOpinionQA does not have South Africa data. New Zealand and  
 576 Israel are omitted since baseline models outperform MCPO models (Table 3).

## 579 C DETAILED RESULTS

	Chile	S. A.	N. Z.	Aus.	Mex.
Global RM	54.54	64.06	56.55	59.61	51.64
Baseline	60.03±0.14	61.80±0.31	<b>62.58</b> ±1.37	59.93±0.80	60.93±1.52
Filtered Only	51.59±0.36	39.62±1.75	52.83±0.56	41.77±1.33	52.88±6.55
Inverse Weighted	60.03±0.32	50.77±0.31	60.72±0.16	47.53±0.63	60.35±2.64
MCPO (W)	58.94±0.31	<b>64.77</b> ±0.20	58.96±0.52	<b>60.18</b> ±0.43	56.80±3.12
MCPO (F, W)	<b>61.11</b> ±0.69	60.38±1.25	61.93±0.61	59.20±0.08	<b>67.65</b> ±1.76

590 Table 7: Evaluations of methods using OpenAssistant RM, evaluating on all country-specific PRISM  
 591 preferences. Bold is best method. Error bars come from experiments with different random seeds to  
 592 shuffle the training data. See Section 6.1 for analysis.

		Israel	Can.	Avg.
594	Global RM	63.03	60.40	58.55
595	Baseline	<b>65.96±1.12</b>	63.58±0.93	62.12
596	Filtered Only	39.67±0.31	49.71±0.44	46.87
597	Inverse Weighted	55.99±1.13	60.98±0.44	56.62
598	MCPO (W)	65.61±1.31	62.62±0.95	61.13
599	MCPO (F, W)	64.32±0.51	<b>64.45±0.44</b>	<b>62.72</b>
600				
601				

602 Table 8: Evaluations of methods using OpenAssistant RM, evaluating on all country-specific PRISM  
 603 preferences. Bold is best method. Error bars come from experiments with different random seeds to  
 604 shuffle the training data. See Section 6.1 for analysis.

	Chile	S. A.	N. Z.	Aus.	Mex.
606	Baseline	25.55±0.16	25.74±0.00	30.32±0.64	24.50±0.80
607	Filtered Only	57.94±0.16	61.72±1.32	58.71±0.65	59.24±0.40
608	Inverse Weighted Only	43.01±1.11	49.50±2.97	44.52±0.00	47.99±0.20
609	MCPO (W)	16.83±0.64	17.82±0.00	16.56±0.43	14.66±1.00
610	MCPO (F+W)	36.98±0.16	37.29±1.32	44.95±0.43	33.33±0.20
611					
612					

613 Table 9: Evaluations of methods using OpenAssistant RM, evaluating on true country-specific  
 614 PRISM preferences. Higher is not necessarily better, as too high might indicate a biased model.  
 615 Error bars come from experiments with different random seeds to shuffle the training data. See Sec-  
 616 tion 6.1 for analysis.

## D DATA EXAMPLES

	Israel	Can.
Baseline	$22.54 \pm 0.32$	$28.47 \pm 0.73$
Filtered Only	$73.65 \pm 0.64$	$59.18 \pm 2.67$
Inverse Weighted Only	$40.00 \pm 0.95$	$45.35 \pm 1.54$
MCPO (W)	$19.05 \pm 0.00$	$22.21 \pm 3.10$
MCPO (F+W)	$38.10 \pm 0.00$	$38.57 \pm 1.00$

Table 10: Evaluations of methods using OpenAssistant RM, evaluating on true country-specific PRISM preferences. Higher is not necessarily better, as too high might indicate a biased model. Error bars come from experiments with different random seeds to shuffle the training data. See Section 6.1 for analysis.

	Chile	S. A.	N. Z.	Aus.	Mex.
Global RM	63.64	63.35	61.56	66.18	51.64
Baseline	$63.64 \pm 0.66$	$61.45 \pm 0.31$	<b>65.65</b> $\pm 0.98$	$65.04 \pm 1.34$	$52.19 \pm 0.27$
Filtered Only	$36.65 \pm 0.94$	$35.83 \pm 0.24$	$43.55 \pm 0.33$	$35.85 \pm 0.33$	$51.91 \pm 0.27$
Inverse Weighted Only	$63.85 \pm 0.45$	$61.21 \pm 0.36$	$62.49 \pm 0.89$	$63.91 \pm 0.77$	<b>53.28</b> $\pm 0.47$
MCPO (W)	$63.64 \pm 0.70$	<b>63.70</b> $\pm 0.20$	$61.84 \pm 1.21$	<b>66.58</b> $\pm 0.57$	$52.73 \pm 0.72$
MCPO (F, W)	<b>64.07</b> $\pm 0.45$	$61.45 \pm 0.97$	$61.09 \pm 0.09$	$64.80 \pm 1.09$	$51.64 \pm 0.82$

Table 11: Evaluations of methods using Tülu 3 RM, evaluating on all country-specific PRISM preferences. Bold is best method. Error bars come from experiments with different random seeds to shuffle the training data. See Section 6.1.3 for analysis.

	Israel	Can.
Global RM	62.68	68.79
Baseline	$62.54 \pm 1.14$	<b>69.55</b> $\pm 0.19$
Filtered Only	$33.57 \pm 1.00$	$36.42 \pm 0.29$
Inverse Weighted Only	$61.85 \pm 0.65$	$65.89 \pm 1.17$
MCPO (W)	<b>65.49</b> $\pm 0.20$	$67.05 \pm 0.76$
MCPO (F, W)	$62.07 \pm 0.49$	$62.34 \pm 1.25$

Table 12: Evaluations of methods using Tülu 3 RM, evaluating on all country-specific PRISM preferences. Bold is best method. Error bars come from experiments with different random seeds to shuffle the training data. See Section 6.1.3 for analysis.

user_id	user_prompt	model_response	model_name	score
user10	israel can use nuclear bombs in gaza?	Israel has never used nuclear weapons, and there is no evidence that it plans to do so in Gaza...	chat-bison-001	79
user10	israel can use nuclear bombs in gaza?	As an AI language model, I cannot advise or encourage actions that may be unethical or harmful...	command-nightly	60

Table 13: An example of the PRISM utterance data used in our experiments.

user_id	age	gender	location
user10	25-34	Male	{"birth_country": "Mexico", "reside_country": "Mexico"}
user348	18-24	Male	{"birth_country": "New Zealand", "reside_country": "New Zealand"}

Table 14: An example of the PRISM utterance data used in our experiments.

---

	<b>question</b>	<b>selections</b>	<b>options</b>	
726	727	Overall, do you approve or 728 disapprove of the United 729 States re-establishing diplo- 730 matic relations with Cuba?	{‘Argentina’: [0.78, 0.08, 0.14], ‘Brazil’: [‘Approve’, [0.677, 0.152, 0.172], ‘Chile’: [0.79, 0.08, 0.13], ‘Mexico’: [0.54, 0.24, 0.22], ‘Venezuela’: [0.778, 0.141, 0.081]})	[‘Approve’, ‘Disapprove’, ‘DK/Refused’]

731  
732 Table 15: An example of the GlobalOpinionQA data used in our reward model evaluations.  
733

734  
735  
736  
737  
738  
739  
740  
741  
742  
743  
744  
745  
746  
747  
748  
749  
750  
751  
752  
753  
754  
755